

REFERENCES

- [1] D. W. Otter, J. R. Medina, and J. K. Kalita, “A survey of the usages of deep learning for natural language processing,” *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 32, no. 2, pp. 604–624, Feb. 2021.
- [2] W. Shi, J. Cao, Q. Zhang, Y. Li, and L. Xu, “Edge computing: Vision and challenges,” *IEEE Internet Things J.*, vol. 3, no. 5, pp. 637–646, Oct. 2016.
- [3] D. Liu, H. Kong, X. Luo, W. Liu, and R. Subramaniam, “Bringing AI to edge: From deep learning’s perspective,” 2020, arXiv:2011.14808.
- [4] M. Chen et al., “Distributed learning in wireless networks: Recent progress and future challenges,” *IEEE J. Sel. Areas Commun.*, vol. 39, no. 12, pp. 3579–3605, Dec. 2021.
- [5] M. Chen et al., “Federated learning of n-gram language models,” in *Proc. 23rd Conf. Comput. Natural Lang. Learn.*, 2019, pp. 121–130.
- [6] M. E. Peters et al., “Deep contextualized word representations,” in *Proc. Annu. Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol.*, 2018, pp. 2227–2237.
- [7] J. D. M. W. C. Kenton and L. K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. Annu. Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol.*, 2019, pp. 4171–4186.
- [8] Y. Liu et al., “RoBERTa: A robustly optimized bert pretraining approach,” 2019, arXiv:1907.11692.
- [9] Z. Lan, M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut, “ALBERT: A lite bert for self-supervised learning of language representations,” 2019, arXiv:1909.11942.
- [10] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT, A distilled version of bert: Smaller, faster, cheaper and lighter,” 2019, arXiv:1910.01108.
- [11] T. Brown et al., “Language models are few-shot learners,” in *Proc. 34th Int. Conf. Neural Inf. Process. Syst.*, 2020, Art. no. 159.
- [12] Z. Zhang, Y. Yang, Y. Dai, L. Qu, and Z. Xu, “When federated learning meets pre-trained language models’ parameter-efficient tuning methods,” 2022, arXiv:2212.10025.
- [13] A. Hard et al., “Federated learning for mobile keyboard prediction,” 2018, arXiv:1811.03604.
- [14] X. Dong, Z. Yu, W. Cao, Y. Shi, and Q. Ma, “A survey on ensemble learning,” *Front. Comput. Sci.*, vol. 14, no. 2, pp. 241–258, 2020.
- [15] Y. Zhou et al., “Mixture-of-experts with expert choice routing,” in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2022, pp. 7103–7114.
- [16] Q. Cao, Y. K. Lal, H. Trivedi, A. Balasubramanian, and N. Balasubramanian, “IrEne: Interpretable energy prediction for transformers,” in *Proc. Annu. Meeting Assoc. Comput. Linguistics*, 2021, pp. 2145–2157.
- [17] A. Vaswani et al., “Attention is all you need,” in *Proc. 31st Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 6000–6010.
- [18] K. Liao, Y. Zhang, X. Ren, Q. Su, X. Sun, and B. He, “A global past future early exit method for accelerating inference of pre-trained language models,” in *Proc. North Amer. Chapter Assoc. Comput. Linguistics*, 2021, pp. 2013–2023.

- [19] P. Auer, N. Cesa-Bianchi, and P. Fischer, “Finite-time analysis of the multiarmed bandit problem,” *Mach. Learn.*, vol. 47, no. 2, pp. 235–256, 2002.
- [20] S. Ramaswamy, R. Mathews, K. Rao, and F. Beaufays, “Federated learning for emoji prediction in a mobile keyboard,” 2019, arXiv:1906.04329.
- [21] M. Chen et al., “Evaluating large language models trained on code,” 2021, arXiv:2107.03374.
- [22] M. Liu, S. Ho, M. Wang, L. Gao, Y. Jin, and H. Zhang, “Federated learning meets natural language processing: A survey,” 2021, arXiv:2107.12603.
- [23] M. Chen, R. Mathews, T. Ouyang, and F. Beaufays, “Federated learning of out-of-vocabulary words,” 2019, arXiv:1903.10635.
- [24] Y. Mansour, M. Mohri, J. Ro, and A. T. Suresh, “Three approaches for personalization with applications to federated learning,” 2020, arXiv:2002.10619.
- [25] S. Ji, S. Pan, G. Long, X. Li, J. Jiang, and Z. Huang, “Learning private neural language modeling with attentive aggregation,” in *Proc. Int. Joint Conf. Neural Netw.*, 2019, pp. 1–8.
- [26] J. Stremmel and A. Singh, “Pretraining federated text models for next word prediction,” in *Proc. Future Inf. Commun. Conf.*, 2021, pp. 477–488.
- [27] C.-J. Wu et al., “Machine learning at facebook: Understanding inference at the edge,” in *Proc. IEEE Int. Symp. High Perform. Comput. Architecture*, 2019, pp. 331–344.
- [28] W. Ren et al., “A survey on collaborative DNN inference for edge intelligence,” 2022, arXiv:2207.07812.